

INSTITUT FÜR INFORMATIK
Computer Vision, Computer Graphics
and Pattern Recognition

Universitätsstr. 1 D-40225 Düsseldorf



Classification of data from the ATLAS experiments

Michael Janschek

Bachelor Thesis

| | |
|--------------------|---|
| Beginn der Arbeit: | 10. Dezember 2015 |
| Abgabe der Arbeit: | 10. März 2016 |
| Gutachter: | Prof. Dr. Stefan Harmeling Prof. Dr. Stefan Conrad |

Erklärung

Hiermit versichere ich, dass ich diese Bachelor Thesis selbstständig verfasst habe. Ich habe dazu keine anderen als die angegebenen Quellen und Hilfsmittel verwendet. Alle Stellen, die aus den Quellen entnommen wurden, sind als solche kenntlich gemacht worden. Diese Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Düsseldorf, den 10. März 2016

Michael Janschek

Abstract

Hier kommt eine ca. einseitige Zusammenfassung der Arbeit rein.

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 1 |
| 1.1 | The Higgs Boson and the ATLAS experiments | 1 |
| 1.2 | The Higgs Boson Machine Learning Challenge | 1 |
| 1.3 | Overview | 1 |
| 2 | Understanding the challenge | 3 |
| 2.1 | The data | 3 |
| 2.2 | The formal problem | 4 |
| 2.3 | The evaluation | 5 |
| 3 | Methods of classification | 7 |
| 3.1 | basic data science methodology | 7 |
| 3.2 | Logistic regression | 7 |
| 3.3 | k-nn classification | 7 |
| 3.4 | The winning methods | 7 |
| 4 | Results on the Kaggle data | 9 |
| 4.1 | k-nn classification | 9 |
| 4.2 | Comparision of all methods | 9 |
| 5 | Discussion | 11 |
| 5.1 | Impact of the Challenge | 12 |
| | References | 13 |
| | List of Figures | 14 |
| | List of Tables | 14 |
| A | The detailed description of the features | 14 |

1 Introduction

This chapter first presents the Higgs Boson Machine Learning Challenge and explains its motivation and goals. It is concluded by an overview of the thesis structure.

1.1 The Higgs Boson and the ATLAS experiments

2013 the Nobel prize in physics acknowledged the discovery of the Higgs Boson. A particle the physicist Peter Higgs predicted to exist, giving mass to other elementary particles.

1.2 The Higgs Boson Machine Learning Challenge

Kaggle is an internet community of data scientists, it hosts several competitions posed by businesses or organizations. Further services involve open datasets, a "Jobs Board" and "Kaggle Rankings", a scoreboard based on performances of community-members in Kaggles competitions.

1.2.1 Motivation

1.2.2 Goal

The goal of the Higgs Boson Machine Learning Challenge is to explore the potential of advanced machine learning methods to improve the discovery significance of the experiment. No knowledge of particle physics is required. Using simulated data with features characterizing events detected by ATLAS, your task is to classify events into "tau tau decay of a Higgs boson" versus "background." [Hig15b]

1.3 Overview

In Chap. 2, we will describe the structure of the challenges dataset [Hig15a] and use simple data-analysis methods, to gain first insight about useful features. We will understand the evaluation metric AMS and related formulas we can use as objective functions for optimizing our classifiers. This will be concluded by deriving the formal problem from the challenges task.

Chap. 3 will use first knowledge about the data to choose simple approaches for classification. After these we will describe several more specific and complex methods.

In Chap. 4 we will use the discussed methods and observe their performance on the challenges data. The thesis closes with a discussion about the approaches and their possible influence on other HEP-applications.

2 Understanding the challenge

In Sec. 2.1, the data will be described and first analysis-methods will be used to gain first knowledge about it. As we continue with sections 2.2 and 2.3, we will formulate a problem to solve and learn about the evaluation the challenge used to rank the participants. Doing so, we rely on [ABCG⁺15], as this paper was created by the challenge-organizers for giving participants a useful documentation.

2.1 The data

All data provided by the challenge- and the "opendata.cern"-dataset [Hig15a] was created by the official ATLAS full detector simulator in a two-part-process. The simulator first reproduces proton-proton collisions, called *events*. Then it tracks these via a virtual model of the ATLAS-detector, the resulting data emulates the statistical properties of the real events. By this procedure it is possible to exactly know if an event is a searched *signal*, or *background*. Signal-events are generated by tau tau decay of the Higgs Boson. Background-events originate from three known processes¹ which produce radiation similar to the signal [ABCG⁺15]. The true *class* of an event is contained in the feature *Label* as "*s*" (signal) and "*b*" (background).

Every event has a feature *Weight*, an artifact by the simulation. Summed, the weights are "an unbiased estimate of the expected number of events falling in the same region during a given fixed time interval. In our case, the weights correspond to the quantity of real data taken during the year 2012" [ABCG⁺15]. This relation causes the weight-mean of signal-events to be considerably smaller than background-weights. We will see, that the challenges evaluation utilizes this to punish incorrectly classifying background as signal.

The features *Weight* and *Label* were originally only provided in the training-dataset. The data used in this thesis is expanded by complete *Weight*-, *Label*-features and the Kaggle-specific features *KaggleSet* and *KaggleWeight*. Last one is a normalization of *Weight* for the total number of signal- or background-events with the same *KaggleSet* feature. This information enables us recreating the challenges original datasets using the opendata.cern-dataset.

The physical features are separated in two types. Features containing so-called primitive data, properties of events explicitly measured by the simulated ATLAS detector, use the preamble "*PRI_*" in their names. The second type is called derived data, which are features that have been computed from primitive features. Their labels use the preamble "*DER_*".

| EventID | DER_mass_MMC | ... | Weight | Label | KaggleSet | KaggleWeight |
|---------|--------------|-----|--------|----------|-----------|--------------|
| 133337 | 1.337 | ... | 0.4 | <i>s</i> | <i>v</i> | 0.04 |

Table 1: Representation of the dataset

¹decay of Z boson, W boson and a pair of top quarks [ABCG⁺15]

All features of the dataset are described more detailed in Appendix A.

One might expect decent physics-knowledge as key in succeeding in the challenge, the top-participants did not use a lot domain-knowledge for feature- or method-selection. One goal of the challenges organization was to set a task for data scientists without any physics-background [ABCG⁺15].

2.1.1 Data visualization

In an first attempt to gain information about our data we use basic methods of data analysis. To evaluate a training sets single feature for classification it is common to plot a histogram of the class-labels, in our case "s" and "b", distribution. A useful way to learn about relations between features is to create *scatter plots* of all two-features-combinations. Using these visualizations, we can identify features with good properties for our task.

Figure 1: Put Histograms here!

Figure 2: Put Scatter plots here!

2.2 The formal problem

We use the formal description of the challenge as it is formulated in [ABCG⁺15].

Let $D = (x_i, y_i, w_i)$, $i \in [1, n]$ be the training sample with n events, where:

- $x_i \in \mathbb{R}^d$ is a d -dimensional vector
- $y_i \in \{b, s\}$ is the label
- and $w_i \in \mathbb{R}^+$ is a non-negative weight.

The sums of signal and background-events

$$S = \sum_{y_i=s} w_i$$

and

$$B = \sum_{y_i=b} w_i$$

represent the *expected total number* of signal and background events, during the time of actual data recording.

Let a function $g : \mathbb{R}^d \rightarrow \{b, s\}$ be a binary classifier, a set $G_s = \{x : g(x) = s\}$ is called the *selection region*. Our task is to create $g(x_i)$ while maximizing the Approximate Median Significance (AMS). We shall submit an *index set* $\hat{G}_s = \{i : g(x_i) = s\}$ of points that g classifies as signal and an *index set* $\hat{G}_b = \{i : g(x_i) \neq s\}$ while $\hat{G}_s \cap \hat{G}_b = \emptyset$.

Figure 3: Flowchart describing the AMS calculation

2.3 The evaluation

The evaluation of a single submission to the challenge is related to the common practice in particle physics to rate a discovery by its statistical significance, in this case

$$Z = \sqrt{2 \left(n \ln \left(\frac{n}{\mu_b} \right) - n + \mu_b \right)} \quad (1)$$

where n is the total number of observed events and μ_b is the expected number of background-events.

Often in particle physics a significance of at least $Z=5$ (a five-sigma effect) is regarded as sufficient to claim a discovery [ABCG⁺15].

By estimating $n = s + b$ and $\mu_b = b$ in Eq. (1), we get the *Approximate Median Significance* (AMS)

$$AMS = \sqrt{2 \left((s + b) \ln \left(1 + \frac{s}{b} \right) - s \right)} \quad (2)$$

which is used by high-energy physicists for optimizing the selection region for stronger discovery significance [ABCG⁺15].

For the challenge, a regularization-term b_{reg} was introduced as an artificial shift to b to decrease variance of the AMS, as this makes it easier to compare the participants if the optimal signal region was small. "The value $b_{reg} = 10$ was determined using preliminary experiments." [ABCG⁺15]

This addition to Eq. (2) makes the final evaluation-formula complete:

$$AMS_2 = \sqrt{2 \left((s + b + b_{reg}) \ln \left(1 + \frac{s}{b + b_{reg}} \right) - s \right)} \quad (3)$$

For simplicity, we will call it just AMS, as Eq. (2) will not have further appearances in this thesis.

2.3.1 The leaderboard

In Kaggle-Challenges, competitors are ranked in a leaderboard, rank 1 is the participant who submitted a solution which achieved the best score on the evaluation used in this challenge, in our case the AMS (3). To prevent participants from simply training algorithms on optimizing the leaderboard-score, Kaggle uses a so-called *public* and *private leaderboard*, which use different data from the test set to generate the AMS. The weights were normalized with respect to the number of events used for computing the public and

Figure 4: Compared rankings of the public and private leaderboard

private score, in our dataset [Hig15a] this information is given by the features *KaggleSet* and *KaggleWeight* so it is possible to compute these scores.

During the challenge, the public leaderboard is visible to any visitor of Kaggle, so participants are able to get an evaluation of their submitted solution and work on better classification for a higher rank. After reaching the *final submission deadline*, the private leaderboard is accessible, which shows the final ranking of the challenge and the differences to the public leaderboard. Only the private rank is relevant for winning the challenge.

2.3.2 Alternative objective functions

For classification, a data scientist wants to train a classifier on an *objective function*. Properties of the AMS (like using the logarithm) make it difficult to use it as objective function, some alternatives were proposed by the challenge-creators [ABCG⁺15] and some challenge-participants via the Kaggle-Forum

- alternative objective function: $\frac{s}{\sqrt{b}}$
 - only valid when $s \ll b$ and $b \gg 1$
- weighted AUC [DMNV14]

3 Methods of classification

- setup
- why these own methods, why are they *simple*?
- why these winning methods
 - XGBoost is awesome, why?
 - resources

3.1 basic data science methodology

- feature-selection/-engineering
- crossvalidation
- objective function

why these methods? resources... 42...

3.2 Logistic regression

- why logreg?
- what is logreg?
- optimizing objective function
- performance
- AMS

3.3 k-nn classification

- why kNN?
- what is kNN?
- optimization
- performance
- AMS
- [PPSS14] are noobs :-P

3.4 The winning methods

resources

3.4.1 Neural networks

1st place <https://github.com/melisgl/higgsml> <https://github.com/melisgl/higgsml/blob/master/doc/model.md> 24gb+ ram minimum: titan
gpu 3rd place <https://www.kaggle.com/c/higgs-boson/forums/t/10481/third-place-model-documentation/55390#post55390>

3.4.2 Regularized greedy forest

<https://github.com/TimSalimans/HiggsML> 2nd place Requires 64gb+ RAM

3.4.3 XGBoost

[CH14]

4 Results on the Kaggle data

4.1 k-nn classification

4.2 Comparision of all methods

5 Discussion

5.1 Impact of the Challenge

Use [CGG⁺14]!

- Ersetzen von AMS durch weighted AUC (papers) [DMNV14]
 - video bei flavour of physics
- XGBoost and Kaggle
- weitere CERN-Challenges

[DMNV14]

References

- [ABCG⁺15] ADAM-BOURDARIOS, Claire ; COWAN, Glen ; GERMAIN, Cécile ; GUYON, Isabelle ; KÉGL, Balázs ; ROUSSEAU, David: *Learning to discover: the Higgs boson machine learning challenge*. <http://www.http://opendata.cern.ch/record/329>, January 2015. – Version 2.3
- [CGG⁺14] COWAN, Glen (Hrsg.) ; GERMAIN, Cécile (Hrsg.) ; GUYON, Isabelle (Hrsg.) ; KÉGL, Balázs (Hrsg.) ; ROUSSEAU, David (Hrsg.) ; CERN, CRNS (Veranst.): *NIPS 2014 Workshop on High-energy Physics and Machine Learning*. 2014
- [CH14] CHEN, Tianqi ; HE, Tong: Higgs Boson Discovery with Boosted Trees. In: *NIPS 2014 Workshop on High-energy Physics and Machine Learning* Bd. 42, 2014 (JMLR: Workshop and Conference Proceedings), S. 69–80. – <http://jmlr.csail.mit.edu/proceedings/papers/v42/diaz14.pdf> , Accessed: 2016-01-18
- [DMNV14] DÍAZ-MORALES, Roberto ; NAVIA-VÁZQUEZ Ángel: Optimization of AMS using Weighted AUC optimized models. In: *NIPS 2014 Workshop on High-energy Physics and Machine Learning* Bd. 42, 2014 (JMLR: Workshop and Conference Proceedings), S. 109–127. – <http://jmlr.csail.mit.edu/proceedings/papers/v42/diaz14.pdf> , Accessed: 2016-01-18
- [Hig15a] *Dataset from the ATLAS Higgs Boson Machine Learning Challenge 2014*. <http://www.http://opendata.cern.ch/record/328>, 2015. – Accessed: 2015-12-10
- [Hig15b] *Higgs Boson Machine Learning Challenge*. <https://www.kaggle.com/c/higgs-boson>, 2015. – Accessed: 2016-01-03
- [PPSS14] PEREZ, Jocelyn ; PONMALAI, Ravi ; SILVER, Alex ; STRACK, Dacoda: *ML2014: Higgs Boson Machine Learning Challenge*. <http://www.http://opendata.cern.ch/record/331>, 2014. – Accessed: 2015-12-10

List of Figures

| | | |
|---|---|---|
| 1 | Put Histograms here! | 4 |
| 2 | Put Scatter plots here! | 4 |
| 3 | Flowchart describing the AMS calculation | 5 |
| 4 | Compared rankings of the public and private leaderboard | 6 |

List of Tables

| | | |
|---|---|---|
| 1 | Representation of the dataset | 3 |
|---|---|---|

A The detailed description of the features